

Compact, Interleaved, Portable: An Open Reference Implementation of the Batched Small-Matrix Linear-Algebra API

A cross-vendor performance study against Intel MKL Compact and Arm Performance Libraries

Ivan Pribec*

August 6, 2026 ▶ [preprint draft — not for circulation]

Abstract

Many scientific kernels do not factor one large matrix but millions of tiny ones, all the same shape: the local systems of meshless approximation, the mapping problems of multi-physics coupling, per-cell chemistry, and per-column implicit solves. The leading vendor answer is the *compact*, or *interleaved-batch*, storage format, in which element (i, j) of V matrices is laid out contiguously so that a single SIMD instruction advances V independent factorizations at once. Intel MKL (Compact) and Arm Performance Libraries (interleave-batch) both ship such kernels — but each is closed, locked to one instruction-set family, and, in MKL’s case, incomplete: it can factor and triangular-solve a batch but offers no supported way to *apply* the orthogonal factor Q . We present `cqr`, an open reference implementation of the compact API written as a *single* vector-typed source that the compiler lowers to SSE, AVX, AVX-512 and, in progress, NEON/SVE. It is a drop-in alternative to `mk1_?geqrf_compact`, it supplies the missing `?ormqr_compact` apply- Q step, and it is intended to interoperate with a vendor’s own compact routines. We describe the format and the two vendor APIs, the portable-SIMD design and its one non-trivial kernel (a branch-free `larfg`), and benchmark the batched QR factorization against MKL Compact on x86 and against ArmPL interleave-batch on Arm. ▶ [one-sentence headline result once measured runs are in.]

1 Introduction

Dense linear algebra is usually told as a story about *one* large matrix: block the operation, reuse the cache, saturate the vector units. A large and growing class of applications tells the opposite story. They must factor or solve a *batch* of many small matrices — orders of a few up to a few hundred — that share a common shape, and no single one of them is large enough to keep a modern core busy. Calling a tuned LAPACK routine once per matrix wastes the machine: the matrices are too small to amortize call overhead or to vectorize internally, and the loop leaves the vector lanes mostly idle.

1.1 Motivating applications

Two applications motivate this work directly; several others share its shape.

* ▶ [confirm author list, affiliation, e-mail, ORCID; add co-authors and acknowledgements.]

Meshless approximation: MLS and RBF-FD. Moving least squares (MLS) and radial-basis-function finite differences (RBF-FD) build differential operators on *scattered* point clouds without a mesh [6, 12]. Around each node one gathers a local stencil of its k nearest neighbours and solves a small dense problem for that node’s differentiation weights: a (weighted) least-squares fit of a polynomial or RBF basis in the MLS case, or a saddle-point RBF-plus-polynomial interpolation system in the RBF-FD case. A discretization with N nodes produces N such systems — routinely 10^6 – 10^8 of them — *all of the same order*, set by the stencil size. Typical stencils of 30–170 points in 2-D and 3-D are exactly the non-power sizes (30, 45, 60, 105, 168) exercised by our benchmark (Section 5). This is the canonical batched-small-matrix workload, and it is the application we care about most.

Data mapping for multi-physics coupling. Partitioned multi-physics couples separately discretized solvers across a shared interface, and their surface meshes rarely match. Coupling libraries such as preCICE [5] therefore *map* field data between non-matching point sets, and radial-basis-function interpolation is a standard choice. Local or partition-of-unity RBF mappings assemble and solve one small dense system per output point or per partition — again many systems of a common size, re-solved every time the interface geometry changes. The batched compact format is a natural fit for the assembly-and-solve inner loop.

Others of the same shape. The pattern recurs. Implicit *chemical kinetics* in reacting-flow solvers integrates a stiff ODE per cell, whose Newton step is a small dense solve of the species Jacobian — one per cell, identically sized. ► [add a kinetics reference (e.g. a GPU/batched chemistry paper).] *PDE line-solvers* — ADI sweeps, implicit vertical transport in atmosphere/ocean columns — produce one small (often banded, sometimes dense) system per grid line. And sequential *convex optimization* (interior-point, SQP, model-predictive control) repeatedly solves small dense KKT systems; the GPL-licensed *batmat* project [11] targets exactly this domain with interleaved batched kernels. ► [add an optimization/MPC reference.] What unites them is not the mathematics but the memory pattern: a large batch of identically shaped small matrices, ripe for cross-matrix vectorization.

1.2 The compact (interleaved) format

The vendor response is a storage format built for this pattern. Rather than store the V matrices of a group back to back, the *compact* (Intel) or *interleaved-batch* (Arm) format *interleaves* them, so that the V copies of element (i, j) sit in V contiguous words. A scalar factorization kernel then lifts almost verbatim to SIMD: replace `double` with a V -wide vector type and every arithmetic operation advances V independent matrices in lockstep (Section 2). Intel MKL exposes this through its Compact BLAS/LAPACK functions [8]; Arm Performance Libraries through its *interleave-batch* functions [3]. Both grew out of the wider batched-BLAS effort [1, 4].

1.3 The gap

Useful as they are, the vendor kernels share three limitations. They are *closed*, so they cannot be studied, ported, or tuned for a new size regime. They are *locked to one ISA family* — MKL Compact to x86, ArmPL to AArch64 — even though the interleaved idea is architecture-neutral. And MKL’s compact LAPACK is *incomplete* for the QR pipeline: it ships `mk1_?geqrf_compact` (factorization) and `mk1_?trsm_compact` (triangular solve) but no `?ormqr_compact`, so there is no supported way to apply the orthogonal factor Q (or Q^T) to a batch — the step that turns a factorization into a least-squares solve.




Figure 1: The compact / interleaved-batch layout for a pack of $V=4$ matrices of order 2 (column-major). The V copies of each element (i, j) occupy V contiguous words, so a single V -wide SIMD load gathers the same element of all V matrices, and a scalar kernel becomes a batched one by widening its scalars to vectors.

1.4 Contributions

This paper presents `cqr`, an open, portable reference implementation of the compact API, and studies its performance against both vendors. Specifically:

- We position the compact/interleaved format and lay the two vendor APIs side by side, making the portability and completeness gaps explicit (Sections 2–3).
- We describe a portable-SIMD implementation from a *single* source — GNU vector types lowered by the compiler to SSE/AVX/AVX-512 (and, in progress, NEON/SVE) — and its one genuinely hard kernel, a branch-free vectorized `larfg` (Section 4).
- We supply the apply- Q step (`cqr_mkl_?ormqr_compact`) that MKL omits, completing the compact QR solve pipeline.
- We benchmark the batched QR factorization against MKL Compact on x86 and ArmPL interleaved-batch on Arm, over the small-size range that the applications above actually use (Sections 5–6).

`cqr` is real-precision (single and double) and unpivoted by design; its scope limits mirror the vendors’ own (Section 7).

2 The compact (interleaved) format

Consider a batch of n_m matrices, all $m \times n$. Group them into $\lceil n_m/V \rceil$ *packs* of V matrices, where V is the SIMD width (for double precision, $V = 2/4/8$ for SSE/AVX/AVX-512). Within a pack, the compact format stores the V values of a given element contiguously; consecutive elements (in the layout’s natural order — column-major here) follow one after another. Figure 1 shows a pack of $V = 4$ matrices of order 2.

This differs from the other common batched convention, the *pointer array* of BBLAS [1], in which each matrix keeps its ordinary dense layout and the batch is an array of pointers. Pointer arrays impose no repacking and suit routines (like batched GEMM) whose per-matrix work is already large enough to vectorize internally; the interleaved layout instead extracts parallelism *across* matrices, which is precisely what wins when each matrix is far too small to

Table 1: The compact/interleaved APIs and `cqr`. `cqr` follows MKL’s calling convention so it can substitute for, or interoperate with, MKL’s own compact routines, while remaining ISA-portable.

	MKL Compact	ArmPL <code>interleave-batch</code>	<code>cqr</code> (this work)
Layout handle	MKL_COMPACT_PACK (opaque V)	explicit inner/outer/batch strides	MKL_COMPACT_PACK <i>or</i> explicit width V
QR factor	<code>mkl_?geqrf_compact</code>	<code>armpl_?geqrf_interleave_batch</code>	<code>cqr_mkl_?geqrf_compact</code>
Apply Q	— (absent)	<code>armpl_?ormqr_interleave_batch</code>	<code>cqr_mkl_?ormqr_compact</code>
Triangular solve	<code>mkl_?trsm_compact</code>	<code>armpl_?trsm_interleave_batch</code>	(use the vendor’s)
Precisions	s,d,c,z	s,d,c,z	s,d (real only)
Target ISA	x86 (SSE/AVX/AVX-512)	AArch64 (NEON/SVE)	any (single vector-typed source)
Source	closed	closed	open (BSD-style)
► [verify exact ArmPL routine names/signatures against the current ArmPL reference.]			

fill a vector register on its own. The cost is a pack/unpack step at the batch boundary (MKL’s `mkl_?gepack_compact/geunpack_compact`); in the target applications the same factored batch is reused enough — across right-hand sides, time steps, or Newton iterations — to amortize it.

3 Two vendor APIs, side by side

Intel MKL Compact. MKL hides the interleave width behind an opaque MKL_COMPACT_PACK token obtained from `mkl_get_format_compact()`, and pairs pack/unpack helpers with a handful of compact BLAS/LAPACK kernels [8]. For QR it provides the factorization and the triangular solve but, as Table 1 records, *no* apply- Q : a caller can obtain the reflectors (H, τ) and can do a triangular solve, but cannot form $Q^T B$ through a supported call. MKL’s compact routines also skip argument checking “for performance reasons” and leave `info` reserved [8], and they document explicit numerical limitations — notably no safe handling of values near underflow/overflow, and no pivoting [9].

Arm Performance Libraries `interleave-batch`. ArmPL exposes the same idea with a more explicit hand: the caller passes inner, outer, and batch strides describing the interleave directly [3], which is more flexible but couples client code to the layout. Crucially for the QR pipeline, ArmPL *does* provide an interleaved apply- Q .

`cqr`. `cqr` adopts MKL’s convention — the opaque MKL_COMPACT_PACK token and the identical argument list — so a call site can swap `mkl_?geqrf_compact` for `cqr_mkl_?geqrf_compact` without change, and `cqr`’s output feeds MKL’s `trsm_compact` unmodified. It also exposes a second, MKL-free C API taking the width V explicitly, for use where MKL is not present. And it fills the apply- Q hole with `cqr_mkl_?ormqr_compact`. The one thing it does not inherit from either vendor is the ISA lock (below).

4 `cqr`: an open, portable reference implementation

One source, every width. `cqr`’s central decision is to write each kernel once, over a scalar type T and an interleave width V , using GNU vector types (`__attribute__((vector_size))`). The

compiler lowers that single source to SSE, AVX, or AVX-512 on x86 — and to NEON or SVE on AArch64 — with no per-ISA intrinsics. A thin dispatcher maps the runtime `MKL_COMPACT_PACK` (or an explicit width) to the compile-time `V` that selects the specialization. This is what lets one code base stand in for two vendors’ closed kernels.

Vectorized unblocked Householder QR. The factorization is the unblocked LAPACK algorithm `geqr2` (`larfg` to build each reflector, `larf` to apply it to the trailing columns [2, 7]), run V matrices at a time. Because the compact layout makes element (i, j) of all V matrices contiguous (Figure 1), the scalar algorithm lifts to SIMD by widening `double` to a V -wide vector: every $+$, $-$, $\times$, $\div$ becomes a lane-wise operation over V independent matrices. The *blocked* algorithm (`geqrf`, with `larft/larfb`) is deliberately *not* used: at these orders the reflector panels are short, and the interleaved batch already saturates the vector units without the extra bookkeeping. Column-major is the tuned path — a matrix column is contiguous in the compact buffer, so the norm reduction, the reflector scaling, and the trailing-column update all walk contiguous pointers, register-blocked a few trailing columns at a time.

The one hard part: a branch-free `larfg`. Scalar `larfg` contains a data-dependent branch (if the below-diagonal norm is zero, $\tau = 0$) and scalar special functions. Across a pack that branch would diverge per lane, so it is rewritten branch-free with a mask and a select, keyed on the LAPACK-correct quantity (the *below-diagonal* norm, so an already-triangular column yields $\tau = 0$ with the diagonal untouched, exactly as `larfg` does):

```
tail = sum_{i>k} A(i,k)^2           // vector reduction below the diagonal
x0   = A(k,k)
norm = sqrt(x0*x0 + tail)           // == hypot(x0, ||tail||)
beta = (x0 >= 0) ? -norm : norm     // -copysign(norm, x0), branch-free
has   = tail > 0                    // lane mask: is there anything to zero?
tau   = has ? (beta - x0)/beta : 0
inv   = has ? 1/(x0 - beta) : 0      // masked select kills 0/0 and 1/0
A(k,k) = has ? beta : x0             // R diagonal (unchanged if no tail)
A(i>k,k) *= inv                     // reflector body (0*inv = 0 when !has)
```

Only `sqrt` needs a helper — a short lane loop that both GCC and Clang lower to a single `vsqrt` — while the comparison yields a native mask and the select is a bit-blend. Every step is $O(1)$ per column; the $O(n^2)$ reductions and $O(n^3)$ trailing update are pure vector arithmetic. The same `has` mask that makes an already-triangular column a no-op also neutralizes the identity matrices that padding puts in the unused lanes of a partial final pack, so the kernel runs unmasked at full width without ever producing a NaN.

The remainder “staircase”. Because work is organized around the width V and a small register-blocking factor, orders n that are not multiples of those quantities leave a partial tile whose cleanup is proportionally more expensive at small n . The sustained rate therefore climbs in a gentle staircase rather than a smooth curve (Figure 5); the effect is most visible at the smallest orders and is exactly why the benchmark includes non-power sizes drawn from real stencils.

Completing the pipeline and its scope. With `cqr_mkl_?ormqr_compact` supplying the apply- Q step, the three routines factor and solve a batched least-squares/linear system end to end — `geqrf_compact` $\rightarrow$ `ormqr_compact('L','T')` $\rightarrow$ `trsm_compact` — the last borrowed from the vendor. Table 2 summarizes coverage. By design `cqr` is real-precision only, performs no argument checking on the MKL-style surface (like MKL), does not pivot (pivoting needs comparisons that do not vectorize across a pack), and does not implement `larfg`’s underflow/overflow rescaling — the same limitation MKL documents for its compact QR [9].

Table 2: **cqr** feature coverage. The tuned contiguous kernel serves the column-major, left-apply solver path; the other layout/side combinations run through a correctness-first stride-generalized kernel over the same math.

Precisions	single (s), double (d); real only
Layouts	column-major (tuned), row-major (stride-generalized)
geqrf	vectorized unblocked geqr2 , branch-free larfg
ormqr	side $\in \{L, R\}$, trans $\in \{N, T\}$ (C folds to T)
potrf	designed (batched Cholesky); not yet implemented
Checking	none on the MKL-style API; LAPACK-style info = $-j$ on the portable C API
Not supported	complex (c,z); column pivoting; underflow/overflow rescaling

5 Experimental methodology

► [This section states the intended protocol; fill the two machine specifications and confirm the details match the runs that produce the data.]

Machines. **x86:** ► [CPU model, base/all-core clock, cores used, L2/L3, MKL version, GCC/Clang version and flags (e.g. `-O3 -march=native`).] **Arm:** ► [AArch64 CPU, SVE vector length, cores, ArmPL version, compiler and flags.]

Workload. Each data point factors a pool of ► [n_m (e.g. 512)] random, well-conditioned square matrices of a fixed order n , in double precision. The size list spans the target range and deliberately mixes powers of two with the non-power stencil sizes 30, 45, 60, 105, 168 (Section 1.1) so the remainder staircase is visible: {8, 16, 24, 30, 32, 45, 48, 60, 64, 96, 105, 128, 168, 170, 256, 384, 500}.

What is timed. The pool is packed into compact form *once*, up front; only the factorization is timed, and the destroyed input is restored (untimed) before each pass. Reported times are the best of ► [r] passes. The two compact paths (**cqr**, vendor) are driven from an outer parallel loop over packs — the intended usage — with the vendor library’s own threading pinned to one thread; the per-matrix LAPACK baseline parallelizes over matrices the same way. This isolates the factorization *kernel*, not pack/unpack or data movement.

Correctness gate. Every timed configuration is checked (untimed) against per-matrix LAPACKE_?geqrf: the unpacked (H, τ) must match elementwise to a tight tolerance, so each benchmark row doubles as an integration test. Reported **relerr** is this quantity.

Reproducibility. The **cqr** x86 numbers come from the repository’s **bench_geqrf_compact** driver; the ArmPL numbers require the AArch64 harness described in Section 9. Scripts, data, and this document build with a single **make** (see the report README).

6 Results

► [All figures below use **synthetic placeholder data** (report/data/*.dat). They exist to validate the plotting pipeline and the narrative shape; replace them with measured runs before drawing any conclusion. See report/data/README.md.]

Throughput on x86 (vs MKL Compact). Figure 2 plots throughput (matrices/s) against order for **cqr**, **mk1_?geqrf_compact**, and the per-matrix LAPACK loop. The expected pattern, and the one the project reports on its AVX-512 test machine, is that both compact paths dominate the per-matrix loop across the small-size range — where cross-matrix SIMD is the




Figure 2: Batched QR factorization throughput on x86 (AVX-512): `cqr` vs MKL Compact vs a per-matrix LAPACK loop, over the target size range. *Placeholder data.*

only way to fill the vector units — and that `cqr` leads MKL at the smallest orders, the two converging as n grows and per-matrix internal vectorization begins to pay for the vendor kernel. Figure 3 shows the same story as speedups over the per-matrix baseline, with the crossover made explicit. ►[state the geometric-mean speedup and the crossover order once measured.]

Throughput on Arm (vs ArmPL interleaved-batch). Figure 4 repeats the comparison on AArch64 against ArmPL’s interleaved-batch `geqrf`. The question here is the *portability tax*: how close does a single vector-typed source come to a vendor kernel hand-tuned for SVE? ►[report the observed gap; note where `cqr` leads (smallest orders) and where ArmPL’s tuning pulls ahead.]

Sustained rate and the staircase. Figure 5 plots `cqr`’s sustained GFLOP/s on both machines. Two features matter: the rate climbs and saturates as n grows (better vector-unit utilization), and it does so in a gentle staircase, dipping at orders that are not multiples of the interleaved width — the remainder effect of Section 4, and the reason the size list includes the odd stencil sizes.

7 Discussion and limitations

The results (once measured) speak to a specific claim: that the interleaved format’s advantage comes from the *layout*, not from any vendor’s closed microkernel, and can therefore be had portably. If `cqr` tracks MKL on x86 and ArmPL on Arm from one source, the format — not the implementation — is doing the work, and there is no reason for it to remain closed or ISA-locked.

The scope limits are deliberate and mirror the vendors’. `cqr` is real-only; complex Hermitian QR is future work. It does not pivot, so it is unsuited to rank-revealing factorizations — but the target least-squares and interpolation problems are set up to be well conditioned. It omits `larfg`’s underflow/overflow rescaling, matching MKL’s documented compact limitation [9]; inputs scaled to the exponent extremes are out of scope. And the row-major and right-apply paths are correctness-first, not separately SIMD-tuned: the tuned path is the column-major, left-apply solver route the applications use.




Figure 3: Speedup over the per-matrix LAPACK baseline on x86. The batched advantage is largest at the smallest orders and narrows toward the parity line as n grows. *Placeholder data.*

8 Related work

The batched-BLAS effort defined the problem and a pointer-array interface standard [1, 4]. The interleaved/compact layout is the vendors’ alternative for the small-matrix regime: Intel’s Compact functions [8] and Arm’s interleave-batch functions [3]. Portable batched kernels also appear in Kokkos Kernels [10], and the `batmat` project [11] provides interleaved batched linear algebra aimed at convex optimization. `cqr`’s contribution relative to these is narrow and specific: an *open* drop-in for the MKL compact QR that is portable across ISAs from one source and that supplies the apply- Q step MKL omits. ► [if targeting a venue, add the closest compact-QR performance papers and any batched-QR-on-GPU comparison for context.]

9 Conclusion and future work

The compact/interleaved format is the right data structure for batches of many small identically-shaped matrices, and it is architecture-neutral even though today’s implementations are closed and ISA-locked. `cqr` shows the format can be served by a single portable, vector-typed source that stands in for `mk1_?geqrf_compact`, completes the pipeline MKL leaves open, and is meant to match both vendors’ kernels in the size range the motivating applications — meshless approximation and multi-physics data mapping — actually use.

Future work: (i) the AArch64/ArmPL benchmark harness needed to turn Figure 4 into measured results — building `cqr` so the vector types lower to NEON/SVE, and a driver that bridges ArmPL’s stride-based interleave to the compact packs; (ii) the batched Cholesky `cqr_mk1_?potrf_compact`, already specified, for the normal-equations form of the MLS/RBF problems; (iii) complex precisions; and (iv) an end-to-end application benchmark (an RBF-FD weight assembly) rather than the kernel microbenchmark alone.

Availability

► [repository URL, licence, and archived release DOI (e.g. Zenodo) for a citable artifact.]




Figure 4: Batched QR factorization throughput on Arm (SVE): `cqr` vs ArmPL interleave-batch vs a per-matrix LAPACK loop. *Placeholder data.*

Acknowledgements

► [funding, compute resources, and any tool/assistance acknowledgements consistent with the project's conventions.]

References

- [1] Ahmad Abdelfattah, Timothy Costa, Jack Dongarra, Mark Gates, Azzam Haidar, Sven Hammarling, Nicholas J. Higham, Jakub Kurzak, Piotr Luszczek, Stanimire Tomov, and Mawussi Zounon. A set of batched basic linear algebra subprograms and LAPACK routines. *ACM Transactions on Mathematical Software*, 47(3), 2021. doi: 10.1145/3431921.
- [2] E. Anderson, Z. Bai, C. Bischof, S. Blackford, J. Demmel, J. Dongarra, J. Du Croz, A. Greenbaum, S. Hammarling, A. McKenney, and D. Sorensen. *LAPACK Users' Guide*. Society for Industrial and Applied Mathematics, Philadelphia, PA, third edition, 1999.
- [3] Arm Limited. Arm performance libraries: Interleave-batch functions. <https://developer.arm.com/documentation/101004/latest/Interleave-batch-functions>, 2026. Accessed 2026.
- [4] Batched BLAS Working Group. Batched BLAS (bblas): a proposed standard interface for batched BLAS operations. <https://icl.utk.edu/bblas/>, 2026. Accessed 2026.
- [5] Hans-Joachim Bungartz, Florian Lindner, Bernhard Gatzhammer, Miriam Mehl, Klaudius Scheufele, Alexander Shukaev, and Benjamin Uekermann. preCICE – a fully parallel library for multi-physics surface coupling. *Computers & Fluids*, 141:250–258, 2016. doi: 10.1016/j.compfluid.2016.04.003.
- [6] Bengt Fornberg and Natasha Flyer. *A Primer on Radial Basis Functions with Applications to the Geosciences*. Society for Industrial and Applied Mathematics, Philadelphia, PA, 2015.
- [7] Gene H. Golub and Charles F. Van Loan. *Matrix Computations*. Johns Hopkins University Press, Baltimore, MD, fourth edition, 2013.




Figure 5: `cqr` sustained rate (GFLOP/s) vs order on x86 and Arm. The wiggle at non-multiple-of- V orders is the SIMD-remainder staircase. *Placeholder data.*

- [8] Intel Corporation. oneMKL developer reference (c): Compact BLAS and LAPACK functions. <https://www.intel.com/content/www/us/en/docs/onemkl/developer-reference-c/2025-2/compact-blas-and-lapack-functions.html>, 2026. Accessed 2026.
- [9] Intel Corporation. oneMKL developer reference (c): Numerical limitations of compact BLAS and compact LAPACK functions. <https://www.intel.com/content/www/us/en/docs/onemkl/developer-reference-c/2025-2/numerical-limits-compact-blas-compact-lapack.html>, 2026. Accessed 2026.
- [10] Kokkos Kernels developers. Kokkos kernels: Batched API. <https://kokkos.org/kokkos-kernels/>, 2026. Accessed 2026.
- [11] Pieter Pas. batmat: batched small-matrix linear algebra. <https://github.com/tttapa/batmat>, 2026. Accessed 2026.
- [12] Holger Wendland. *Scattered Data Approximation*. Cambridge University Press, 2004.